

April 2, 2023

# 1 Experiment 1 - Train a Machine Learning Model to Detect Loan Fraud

## 1.1 Import Libraries

```
[ ]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

%matplotlib inline
```

## 1.2 Extract Data

```
[ ]: train = pd.read_csv('data/train.csv')
train.head()
```

```
[ ]:      Loan_ID Gender Married Dependents      Education Self_Employed \
0  LP001002   Male      No           0      Graduate           No
1  LP001003   Male     Yes           1      Graduate           No
2  LP001005   Male     Yes           0      Graduate           Yes
3  LP001006   Male     Yes           0  Not Graduate           No
4  LP001008   Male     No           0      Graduate           No

      ApplicantIncome  CoapplicantIncome  LoanAmount  Loan_Amount_Term \
0              5849              0.0          NaN          360.0
1              4583             1508.0          128.0          360.0
2              3000              0.0           66.0          360.0
3              2583             2358.0          120.0          360.0
4              6000              0.0          141.0          360.0

      Credit_History  Property_Area  Loan_Status
0              1.0          Urban            Y
1              1.0          Rural            N
2              1.0          Urban            Y
3              1.0          Urban            Y
4              1.0          Urban            Y
```

```
[ ]: test = pd.read_csv('data/test.csv')
test.head()
```

```
[ ]:      Loan_ID Gender Married Dependents      Education Self_Employed \
0  LP001015   Male     Yes         0      Graduate           No
1  LP001022   Male     Yes         1      Graduate           No
2  LP001031   Male     Yes         2      Graduate           No
3  LP001035   Male     Yes         2      Graduate           No
4  LP001051   Male     No          0  Not Graduate           No

      ApplicantIncome  CoapplicantIncome  LoanAmount  Loan_Amount_Term \
0                5720                 0         110.0             360.0
1                3076             1500         126.0             360.0
2                5000             1800         208.0             360.0
3                2340             2546         100.0             360.0
4                3276                 0          78.0             360.0

      Credit_History  Property_Area
0                1.0           Urban
1                1.0           Urban
2                1.0           Urban
3                NaN           Urban
4                1.0           Urban
```

### 1.3 Explore the Dataset

```
[ ]: train.shape
```

```
[ ]: (614, 13)
```

```
[ ]: test.shape
```

```
[ ]: (367, 12)
```

```
[ ]: list(train.columns)
```

```
[ ]: ['Loan_ID',
      'Gender',
      'Married',
      'Dependents',
      'Education',
      'Self_Employed',
      'ApplicantIncome',
      'CoapplicantIncome',
      'LoanAmount',
      'Loan_Amount_Term',
      'Credit_History',
```

```
'Property_Area',  
'Loan_Status']
```

```
[ ]: list(test.columns)
```

```
[ ]: ['Loan_ID',  
      'Gender',  
      'Married',  
      'Dependents',  
      'Education',  
      'Self_Employed',  
      'ApplicantIncome',  
      'CoapplicantIncome',  
      'LoanAmount',  
      'Loan_Amount_Term',  
      'Credit_History',  
      'Property_Area']
```

```
[ ]: train.dtypes
```

```
[ ]: Loan_ID          object  
      Gender          object  
      Married         object  
      Dependents      object  
      Education       object  
      Self_Employed   object  
      ApplicantIncome  int64  
      CoapplicantIncome float64  
      LoanAmount       float64  
      Loan_Amount_Term float64  
      Credit_History   float64  
      Property_Area    object  
      Loan_Status      object  
      dtype: object
```

```
[ ]: train['Loan_Status'].value_counts()
```

```
[ ]: Y    422  
      N    192  
      Name: Loan_Status, dtype: int64
```


```
[ ]: train['Loan_Status'].value_counts(normalize=True)
```

```
[ ]: Y    0.687296  
      N    0.312704  
      Name: Loan_Status, dtype: float64
```

### 1.3.1 Check whether the dataset is balanced or not

```
[ ]: train['Loan_Status'].value_counts().plot.bar()
```

```
[ ]: <AxesSubplot: >
```



```
[ ]: train['Dependents'].replace('3+', 3,inplace=True)
test['Dependents'].replace('3+', 3,inplace=True)

train['Loan_Status'].replace('N', 0,inplace=True)
train['Loan_Status'].replace('Y', 1,inplace=True)
```


### 1.3.2 Plotting a correlation heatmap between the features of the dataset

```
[ ]: matrix = train.corr()
f, ax = plt.subplots(figsize=(9,6))
sns.heatmap(matrix, vmax=.8, cmap="BuPu", annot = True)
```

/tmp/ipykernel\_19879/130472854.py:1: FutureWarning: The default value of numeric\_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric\_only to silence this warning.

```
matrix = train.corr()
```

```
[ ]: <AxesSubplot: >
```



### 1.3.3 Replacing null values with the modes of the respective feature

```
[ ]: train['Gender'].fillna(train['Gender'].mode()[0], inplace=True)
train['Married'].fillna(train['Married'].mode()[0], inplace=True)
train['Dependents'].fillna(train['Dependents'].mode()[0], inplace=True)
train['Self_Employed'].fillna(train['Self_Employed'].mode()[0], inplace=True)
train['Credit_History'].fillna(train['Credit_History'].mode()[0], inplace=True)

train['Loan_Amount_Term'].fillna(train['Loan_Amount_Term'].mode()[0],
    ↪inplace=True)

train['LoanAmount'].fillna(train['LoanAmount'].median(), inplace=True)
```

```
[ ]: train.isnull().sum()
```

```
[ ]: Loan_ID          0
      Gender          0
```

```

Married          0
Dependents       0
Education        0
Self_Employed   0
ApplicantIncome  0
CoapplicantIncome 0
LoanAmount       0
Loan_Amount_Term 0
Credit_History  0
Property_Area    0
Loan_Status      0
dtype: int64

```

```

[ ]: test['Gender'].fillna(train['Gender'].mode()[0], inplace=True)
test['Married'].fillna(train['Married'].mode()[0], inplace=True)
test['Dependents'].fillna(train['Dependents'].mode()[0], inplace=True)
test['Self_Employed'].fillna(train['Self_Employed'].mode()[0], inplace=True)
test['Credit_History'].fillna(train['Credit_History'].mode()[0], inplace=True)
test['Loan_Amount_Term'].fillna(train['Loan_Amount_Term'].mode()[0],
    ↪inplace=True)
test['LoanAmount'].fillna(train['LoanAmount'].median(), inplace=True)

```

```

[ ]: test.isnull().sum()

```

```

[ ]: Loan_ID      0
Gender           0
Married          0
Dependents       0
Education        0
Self_Employed   0
ApplicantIncome  0
CoapplicantIncome 0
LoanAmount       0
Loan_Amount_Term 0
Credit_History  0
Property_Area    0
dtype: int64

```

```

[ ]: loanID = train['Loan_ID']

train = train.drop('Loan_ID',axis=1)
test = test.drop('Loan_ID',axis=1)

```

```

[ ]: X = train.drop('Loan_Status', 1)
y = train.Loan_Status

```

/tmp/ipykernel\_19879/1314552169.py:1: FutureWarning: In a future version of

pandas all arguments of DataFrame.drop except for the argument 'labels' will be keyword-only.

```
X = train.drop('Loan_Status', 1)
```

## 1.4 Train-test Split

```
[ ]: X = pd.get_dummies(X)
train = pd.get_dummies(train)
test = pd.get_dummies(test)
```

```
[ ]: from sklearn.model_selection import train_test_split
x_train, x_cv, y_train, y_cv = train_test_split(X,y, test_size=0.3)
```

## 1.5 Fitting the Data to a Logistic Regression Model

```
[ ]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

model = LogisticRegression()
model.fit(x_train, y_train)
```

```
[ ]: LogisticRegression()
```

```
[ ]: x_cv
```

```
[ ]:
ApplicantIncome  CoapplicantIncome  LoanAmount  Loan_Amount_Term  \
351              8750              4167.0        308.0             360.0
399              1500              1800.0        103.0             360.0
337              2500              4600.0        176.0             360.0
456              4301               0.0        118.0             360.0
576              3087              2210.0        136.0             360.0
..              ...              ...          ...             ...
380              3333              2500.0        128.0             360.0
92               3273              1820.0         81.0             360.0
317              2058              2134.0         88.0             360.0
240              5819              5000.0        120.0             360.0
60               2500              3796.0        120.0             360.0

Credit_History  Gender_Female  Gender_Male  Married_No  Married_Yes  \
351             1.0            0            1            1            0
399             0.0            1            0            1            0
337             1.0            0            1            0            1
456             1.0            0            1            0            1
576             0.0            0            1            0            1
..             ...            ...          ...          ...          ...
380             1.0            0            1            0            1
92              1.0            0            1            0            1
```

317	1.0	0	1	0	1
240	1.0	0	1	0	1
60	1.0	0	1	0	1

	Dependents_3	Dependents_0	Dependents_1	Dependents_2	\
351	0	1	0	0	
399	0	1	0	0	
337	0	0	0	1	
456	0	1	0	0	
576	0	1	0	0	
..	...	...	...	...	
380	0	1	0	0	
92	0	0	0	1	
317	0	1	0	0	
240	0	0	0	1	
60	0	1	0	0	

	Education_Graduate	Education_Not Graduate	Self_Employed_No	\
351	1	0	1	
399	1	0	1	
337	1	0	0	
456	1	0	1	
576	1	0	1	
..	...	...	...	
380	1	0	1	
92	0	1	1	
317	1	0	1	
240	1	0	1	
60	1	0	1	

	Self_Employed_Yes	Property_Area_Rural	Property_Area_Semiurban	\
351	0	1	0	
399	0	0	1	
337	1	1	0	
456	0	0	0	
576	0	0	1	
..	...	...	...	
380	0	0	1	
92	0	0	0	
317	0	0	0	
240	0	1	0	
60	0	0	0	

	Property_Area_Urban
351	0
399	0
337	0



```

456          1
576          0
..          ...
380          0
92           1
317          1
240          0
60           1

```

[185 rows x 20 columns]

## 1.6 Accuracy of our logistic regression model

```
[ ]: pred_cv = model.predict(x_cv)

accuracy_score(y_cv, pred_cv)
```

```
[ ]: 0.654054054054054
```

```
[ ]: test
```

```
[ ]:
      ApplicantIncome  CoapplicantIncome  LoanAmount  Loan_Amount_Term  \
0          5720          0          110.0          360.0
1          3076          1500          126.0          360.0
2          5000          1800          208.0          360.0
3          2340          2546          100.0          360.0
4          3276          0           78.0          360.0
..          ...          ...          ...          ...
362         4009          1777          113.0          360.0
363         4158          709          115.0          360.0
364         3250          1993          126.0          360.0
365         5000          2393          158.0          360.0
366         9200          0           98.0          180.0

      Credit_History  Gender_Female  Gender_Male  Married_No  Married_Yes  \
0          1.0          0          1          0          1
1          1.0          0          1          0          1
2          1.0          0          1          0          1
3          1.0          0          1          0          1
4          1.0          0          1          1          0
..          ...          ...          ...          ...
362         1.0          0          1          0          1
363         1.0          0          1          0          1
364         1.0          0          1          1          0
365         1.0          0          1          0          1
366         1.0          0          1          1          0

```

	Dependents_3	Dependents_0	Dependents_1	Dependents_2	\
0	0	1	0	0	
1	0	0	1	0	
2	0	0	0	1	
3	0	0	0	1	
4	0	1	0	0	
..	...	...	...	...	
362	1	0	0	0	
363	0	1	0	0	
364	0	1	0	0	
365	0	1	0	0	
366	0	1	0	0	

	Education_Graduate	Education_Not Graduate	Self_Employed_No	\
0	1	0	1	
1	1	0	1	
2	1	0	1	
3	1	0	1	
4	0	1	1	
..	...	...	...	
362	0	1	0	
363	1	0	1	
364	1	0	1	
365	1	0	1	
366	1	0	0	

	Self_Employed_Yes	Property_Area_Rural	Property_Area_Semiurban	\
0	0	0	0	
1	0	0	0	
2	0	0	0	
3	0	0	0	
4	0	0	0	
..	...	...	...	
362	1	0	0	
363	0	0	0	
364	0	0	1	
365	0	1	0	
366	1	1	0	

	Property_Area_Urban
0	1
1	1
2	1
3	1
4	1
..	...
362	1

```

363          1
364          0
365          0
366          0

```

```
[367 rows x 20 columns]
```

```
[ ]: pred_test = model.predict(test)
```

```
[ ]: prediction = pd.read_csv('data/predicted.csv')
prediction.head()
```

```
[ ]:      Loan_ID Loan_Status
0  LP001015          N
1  LP001022          N
2  LP001031          N
3  LP001035          N
4  LP001051          N

```

### 1.6.1 Creating a dataframe for the predicted values from our model

```
[ ]: prediction['Loan_Status'] = pred_test

prediction['Loan_Status'].replace(0, 'N', inplace=True)
prediction['Loan_Status'].replace(1, 'Y', inplace=True)

predictions = pd.DataFrame(prediction, columns=['Loan_ID', 'Loan_Status'])
predictions

```

```
[ ]:      Loan_ID Loan_Status
0  LP001015          Y
1  LP001022          Y
2  LP001031          Y
3  LP001035          Y
4  LP001051          Y
..      ...          ...
362 LP002971          Y
363 LP002975          Y
364 LP002980          Y
365 LP002986          Y
366 LP002989          Y

```

```
[367 rows x 2 columns]
```

### 1.6.2 Writing that prediction dataframe to a .csv file

```
[ ]: predictions.to_csv('data/logistic.csv')
```